

Weekly Research Update: Integrating Habit Lab Insights, Clustering, and Simulation into the Nudge Framework

Habit Lab and the Transtheoretical Model (TTM)

The **HABITS Lab** at UMBC (University of Maryland, Baltimore County) is led by Dr. Carlo DiClemente (co-developer of the Transtheoretical Model) and focuses on understanding health behavior change in areas like smoking cessation, addiction, exercise, and diet ¹. A core foundation of the lab's approach is the **Transtheoretical Model (TTM)** of behavior change. TTM is an *integrative, biopsychosocial model* that assesses an individual's **readiness to act** on a new, healthier behavior and provides strategies (processes of change) to guide them through behavior change ². In TTM, change is viewed as a progression through **Stages of Change** ³:

- **Precontemplation:** Not yet considering change or unaware of the problem ⁴.
- **Contemplation:** Aware of the need for change and considering action, but ambivalent ⁵.
- **Preparation:** Intending to take action soon and perhaps taking small steps ⁶.
- **Action:** Actively engaging in the new behavior (within recent months) ⁷.
- **Maintenance:** Sustaining the behavior change over time and working to prevent relapse ⁸.

TTM emphasizes that change is not instantaneous – people often **recycle through stages** or regress before progressing again ⁹. This perspective complements our framework's *Behavioral Readiness Assessment* step. In our pseudocode framework, we used a COM-B model (evaluating *Capability, Opportunity, and Motivation*) to assess a patient's readiness/barriers. Integrating TTM insights can **enhance this assessment**. For example, identifying a patient's stage of change can inform which **nudge strategies** are appropriate. A patient in *Precontemplation* (perhaps low motivation and low awareness) might need educational nudges to build capability and awareness, whereas a patient in *Preparation* might benefit more from prompts and planning tools to translate motivation into action. The Habit Lab's work underscores tailoring interventions to an individual's stage and readiness, rather than a one-size-fits-all approach ¹⁰. In fact, TTM has been called "*the dominant model of health behaviour change*" for its broad application to various behaviors ¹¹, and it aligns with our framework's goal of personalized support.

Key takeaway: Incorporating Habit Lab's TTM perspective into our framework means explicitly considering *where* the patient is in their change journey. This could be an added sub-step in the Behavioral Readiness Assessment (e.g., determining if the patient is not ready, getting ready, or taking action) and then choosing nudges that meet them at that stage. By doing so, we ensure the intervention is **stage-matched** – for instance, providing motivational messages to those not yet motivated, versus action-oriented reminders for those already intending to change.

Integrating Clustering (Zane's Work) for Personalization

Our work is directly connected to Zane's preliminary research on **clustering patient data**. The idea is to group patients into clusters (segments) with similar profiles or behavior patterns, so that our framework can tailor interventions to each group. *Why clustering?* Research shows that one-size-fits-all interventions are often less effective than personalized ones, and clustering offers a middle ground between treating everyone the same and treating each person as completely unique ¹⁰. By identifying **homogeneous subgroups**, we can design nudge strategies that target the specific needs or barriers common to each subgroup ¹². This can greatly **speed up learning and optimization** of interventions, as demonstrated in a recent study where a cluster-based reinforcement learning approach outperformed both individual-level and population-level personalization in finding effective health intervention policies ¹² ¹⁰.

In the context of our project (atrial fibrillation patient self-management), clustering can be extremely valuable. A recent exploratory study clustered **AF (Atrial Fibrillation) patients** based on 11 clinical and personal variables (including age, symptom severity, health knowledge, stress levels, quality of life, etc.) ¹³. They found clear groupings – for example, some patients were “doing well” on all measures, while others were struggling with high stress or low satisfaction ¹⁴. Notably, among the patients *not managing well*, there were **three distinct subgroups** characterized as “**stressed and discontented**,” “**struggling and dissatisfied**,” and “**satisfied and complacent**,” each with different challenges ¹⁴. The authors conclude that AF patients can indeed be classified into distinct natural groupings that **vary in clinically important ways**, and that these subgroups “*may benefit from tailored approaches to AF management and support.*” ¹⁵. In other words, **different clusters might require different nudges** ¹⁵ – e.g., a “stressed and discontented” group might need stress-reduction and motivational support, whereas a “complacent” group might need education about risks to prompt action.

For our framework, Zane's clustering results could be **ingested at multiple points**. One approach is to use cluster membership as part of the input to the nudge strategy selection. For instance, after assessing an individual patient, we might detect they belong to a cluster characterized by *low motivation and low knowledge*. Knowing this, the framework can automatically pick a strategy emphasizing motivational and educational nudges (targeting the “Motivation” and “Capability” components from COM-B). Clusters can also inform the **content personalization** – e.g., patients in a “high stress” cluster might receive empathy-driven messages acknowledging stress, while those in a “complacent” cluster might receive more urgent, risk-focused messages to spur action. Essentially, clustering provides a data-driven way to define **personas or user segments** within our system. This bridges Zane's data-driven insights with our theoretical framework: instead of treating every patient as unique (which can be noisy and resource-intensive), we treat cluster identity as a key contextual variable for tailoring the intervention.

Key takeaway: Leveraging clustering enhances personalization by categorizing patients into **actionable segments**. Our framework should incorporate an initial step or module to **map a patient to a cluster** (based on their profile data and behavior history). That cluster can then guide the subsequent steps – from identifying the primary barrier (since each cluster may have a dominant barrier) to choosing the nudge type most effective for that group. This approach is supported by literature showing improved outcomes when interventions are tailored to group profiles ¹⁵, and it aligns with modern trends of using data analytics (like unsupervised learning) to boost health intervention efficacy.

Using Simulation Models to Test and Refine the Framework

With the framework becoming more detailed (incorporating behavioral models and cluster-based personalization), an important next step is to **simulate how it would perform** before deploying it in a real-world setting. Simulation models can act as a *safe sandbox* to experiment with our Just-In-Time Adaptive Intervention logic and ensure the strategies work as intended. In fact, the concept of “**digital twins**” is emerging in the mHealth field – creating virtual representations of patient subpopulations to test interventions. Researchers have introduced *JITAI-Twins*, which are essentially **digital twin simulations of subpopulations** used to **virtually simulate the outcomes** of different intervention design decisions ¹⁶. By doing this, one can tweak the nudge timing, content, or algorithm parameters *in silico* and observe potential results, without risking patient health or engagement in the process ¹⁷ ¹⁶. The simulation can be continually updated with real-world data (a feedback loop) to improve the model’s accuracy over time ¹⁸.

In our project’s context, we could develop a **simulation environment** where synthetic patients (or “virtual patients”) go through the entire pseudocode pipeline. We already have a placeholder for `synthetic_patient_input` in our pseudocode, which implies we can generate mock patient profiles. Using these, we can simulate scenarios such as: a patient misses a dose and triggers an immediate nudge, or a patient’s motivation score gradually improves over time with successive nudges. The simulation would allow us to answer questions like: *Are the nudges being scheduled at appropriate times? Do patients in the simulation respond (or at least the logic records a response) as expected?* If we integrate reinforcement learning in the future, a simulator is critical for training the algorithm on how to optimize nudges. For example, the cluster-based RL study mentioned earlier actually **developed a simulator** to evaluate how well their policies perform over 100 days of simulated user behavior ¹⁹. Such a simulator can model user dynamics (e.g., probability of taking medication when nudged vs not nudged, likelihood of reading educational content, etc.), and we can adjust our strategies based on the simulated outcomes.

By using simulation, we also ensure that our **Just-In-Time Adaptive Intervention (JITAI)** logic (step 6 in pseudocode) is refined. We can simulate edge cases: What if a patient consistently ignores prompts? Does the system escalate or change strategy appropriately? What if two different types of barriers occur simultaneously? Simulation can reveal if our framework needs additional rules or flexibility. In summary, simulation models act as a **proof-of-concept testbed** for the framework – helping to optimize the design (much like A/B testing different strategies virtually) and giving us confidence before actual deployment. This approach of “*virtual trial runs*” is aligned with current best practices in digital health intervention development, where iterative refinement via modeling and small pilot tests precedes wide rollout ²⁰ ¹⁶.

Key takeaway: Incorporating a simulation phase in our project will allow us to **test and fine-tune the nudge framework** under various conditions. It reduces risk and guides configuration choices (e.g., timing algorithms or content variations) by predicting outcomes using virtual patients. Going forward, we might create a “digital twin” of our target patient population – using the data we have (or expect to collect) to simulate how our interventions play out and continuously improve the strategy before and between real-world trials ²¹.

Recommendation: Extend the Existing Paper vs. Create a New Document

Given the above research and developments, a logical question is how to document this week's work in relation to our existing research paper (from October 19). The new insights on Habit Lab/TTM, clustering, and simulation **directly extend the framework** we already outlined. They represent deeper specificity and future implementation steps building on the foundation. Therefore, it would be most coherent to **add one or two new sections to the existing paper** rather than starting a separate paper from scratch. By adding sections, we maintain a single, comprehensive narrative that shows progression from the conceptual model and pseudocode (initial paper) into these new considerations (this week's additions). For example, we could add:

- *"Incorporating Behavioral Models (TTM) and Habit Lab Insights"* – discussing how COM-B and TTM can be used together, as we did above.
- *"Personalization via Clustering and Simulation"* – covering how Zane's clustering work feeds into the framework and proposing a simulation approach to validate the system.

Adding these sections will clearly tie the new content to each step of the framework already described, making the overall paper more robust and closer to a grant-ready proposal. It also shows the supervisor that we are building on the existing work systematically (which is often desirable in a research narrative). A separate paper or notes might lead to redundancy or fragmentation unless the scope had shifted dramatically (which it hasn't – we are enhancing the same project).

In summary, **integrating this week's work as new sections** in the current research paper is the recommended approach. This will result in a single document that captures everything from the high-level theory to the technical implementation ideas, strengthened by the new findings (Habit Lab's theoretical framework and cluster/simulation methods). Such a document could serve as a strong basis if you proceed to write a grant proposal, as it demonstrates both the theoretical grounding and the data-driven, technical plan for your intervention framework.

1 About the Lab – HABITS Lab – UMBC

<https://habitslab.umbc.edu/about-us/>

2 11 Transtheoretical model - Wikipedia

https://en.wikipedia.org/wiki/Transtheoretical_model

3 4 5 6 7 8 9 About the TTM – HABITS Lab – UMBC

<https://habitslab.umbc.edu/the-model/>

10 12 19 research.vu.nl

<https://research.vu.nl/files/246409943/>

[Personalization_of_Health_Interventions_Using_Cluster-Based_Reinforcement_Learning.pdf](#)

13 14 15 Managing With Atrial Fibrillation: An Exploratory Model-Based Cluster Analysis of Clinical and Personal Patient Characteristics - PubMed

<https://pubmed.ncbi.nlm.nih.gov/38020332/>

16 17 18 20 21 Publication: Digital Twins for Just-in-Time Adaptive Interventions (JITAI-Twins): A Framework for Optimizing and Continually Improving JITAIs – Massachusetts AI & Technology Center
<https://massaitc.org/2025/05/09/publication-digital-twins-for-just-in-time-adaptive-interventions-jitai-twins-a-framework-for-optimizing-and-continually-improving-jitais/>